

Automating Software Traceability Link Recovery and Maintenance using Word Embeddings within a Shallow Neural Network

Marlan McInnes-Taylor

Faculty Mentor: Dr. Chris Mills

Florida State University Department of Computer Science

1017 Academic Way, Tallahassee, FL 32304

mcinnest@cs.fsu.edu, crmills@fsu.edu

1 Abstract

In addition to code, software systems contain a multitude of files documenting features, known issues, legal requirements, etc. Software traceability is the ability to link related documents from these various sets through a process called *traceability link recovery*. While previous research has shown that established software traceability improves the quality of projects and makes them easier to maintain, establishing software traceability is often a manual, arduous, and error prone task. This research explores automating the process of software traceability link recovery using established techniques in machine learning. The results demonstrate that even with minimally tuned hyperparameters a shallow neural network can effectively predict which text-based artifacts within a software system are related to one another.

2 Introduction

Previous studies have shown that access to traceability information explaining relationships between software artifacts in a system leads to better software quality and lower bug counts. This is largely credited to such information improving fundamental software engineering tasks such as program comprehension, concept and bug localization, and defect prediction among many others [5, 10, 13]. Unfortunately, the acquisition of traceability information is difficult and typically an afterthought during system construction. Further, when using a manual process to establish traceability, hundreds or thousands of man hours are spent inferring relationships between artifacts that are constantly in flux as the system changes during development and maintenance [8, 14, 2]. Therefore, even if a firm is able to establish traceability for a mission critical system, the effectiveness of that information is potentially temporary. To address this situation, previous studies have attempted to either completely or partially automate the process of inferring traceability links between system components [4]. In this work, we continue that research agenda by using neural networks to model links between text-based software artifacts and predict which of those links are valid – meaning that the link exists between two related artifacts within the software system. Preliminary results of using this technique are promising, with higher recall and precision than its contemporaries [12, 11].

3 Materials and Methods

The training and test data were comprised of the Albergate, Eanci, eTour, iTrust, Modis, and SMOS datasets. All code was written in Python.

3.1 Data Cleaning

All document text was first tokenized using NLTK's [3] punkt tokenizer. The documents were cleaned by removing stopwords. NLTK's built in stopwords list was used as a basis, with a collection of Java and C

stopwords appended to it. The documents were further cleaned by removing whitespace, punctuation, and purely numeric tokens.

3.2 Metadocument Generation

Metadocuments were created using the cleaned data from each dataset. Every metadocument was comprised of text data associated with a use case (UC) and a code class (CC), and metadocuments were created for all possible combinations of a dataset’s UCs and CCs. Additionally, the text of each metadocument was randomly shuffled using the Numpy’s[7] shuffle method, in order to reduce potential bias from token arrangement when setting the sequence length. Finally, each metadocument was classified as either a valid link between two related documents or an invalid link between two unrelated documents, as specified in each dataset’s oracle file.

3.3 Dataset Balancing

Class imbalance was present in all datasets due to the large number of possible UC and CC combinations. As one might expect, for a given system, there are far fewer metadocuments representing valid links than invalid ones. To prevent class imbalance from impacting model performance, each dataset was balanced using Imbalanced-Learn’s[9] SMOTE implementation. Due to the SMOTE parameters, all metadocuments within the set were pruned to a uniform length before synthetic metadocuments were created. The average length of all metadocuments within the set was used for pruning. Synthetic metadocument generation resulted in an equal number of valid and invalid links within each dataset.

3.4 Model Design and Training

A shallow Tensorflow[1] model was implemented to perform metadocument classification. The model’s first two layers performed text vectorization and transformed the vectors using word embeddings. The vectors were then pooled before being passed into a 16 node dense layer followed by the output layer. Initial positive results on the SMOS dataset led to 50 trials of 10 fold cross validation with 15 epochs per fold on each dataset. Shuffled stratification via scikit-learn’s[6] StratifiedShuffleSplit implementation was used to sample the dataset in each fold.

4 Results

Dataset	Loss	Accuracy	Recall	Precision	F1Score
Albergate	0.28	0.95	0.91	0.98	0.94
Eanci	0.11	0.96	0.94	0.98	0.96
Etour	0.13	0.96	0.93	0.99	0.96
iTrust	0.17	0.97	0.95	0.99	0.97
Modis	0.29	0.93	0.94	0.92	0.93
SMOS	0.33	0.87	0.75	0.98	0.85

5 Conclusions and Future Work

Given the strong precision and recall results, it is clear that word embeddings used within a shallow network can successfully perform metadocument link classification within a defined training and testing environment. Future work will focus on improving this model by further optimizing its hyperparameters. To make the model more appealing to industrial partners, we plan to determine ways to minimize the start-up cost associated with employing our model to establish traceability information in a greenfield system. Initially, we plan to perform a series of experiments to identify the minimum data requirements for a given level of performance while employing supportive techniques such as Active Learning. Further, we will focus our efforts on using filtered data from disparate software systems to create a generalizable model for initial traceability link recovery with limited or absent within-project data.

References

- [1] Martín Abadi, Paul Barham, Jianmin Chen, Zhifeng Chen, Andy Davis, Jeffrey Dean, Matthieu Devin, Sanjay Ghemawat, Geoffrey Irving, Michael Isard, et al. Tensorflow: A system for large-scale machine learning. In *12th {USENIX} Symposium on Operating Systems Design and Implementation ({OSDI} 16)*, pages 265–283, 2016.
- [2] Giuliano Antoniol, C Casazza, and Aniello Cimitile. Traceability recovery by modeling programmer behavior. In *Proceedings of the Seventh Working Conference on Reverse Engineering*, pages 240–247. IEEE, 2000.
- [3] Steven Bird, Ewan Klein, and Edward Loper. *Natural Language Processing with Python*. O’Reilly Media, 2009.
- [4] Markus Borg, Per Runeson, and Anders Ardö. Recovering from a decade: A systematic mapping of information retrieval approaches to software traceability. *Empirical Software Engineering*, 19(6):1565–1616, 2014.
- [5] Elke Bouillon, Patrick Mäder, and Ilka Philippow. A survey on usage scenarios for requirements traceability in practice. In *International Working Conference on Requirements Engineering: Foundation for Software Quality*, pages 158–173. Springer, 2013.
- [6] Lars Buitinck, Gilles Louppe, Mathieu Blondel, Fabian Pedregosa, Andreas Mueller, Olivier Grisel, Vlad Niculae, Peter Prettenhofer, Alexandre Gramfort, Jaques Grobler, Robert Layton, Jake VanderPlas, Arnaud Joly, Brian Holt, and Gaël Varoquaux. API design for machine learning software: experiences from the scikit-learn project. In *ECML PKDD Workshop: Languages for Data Mining and Machine Learning*, pages 108–122, 2013.
- [7] Charles R. Harris, K. Jarrod Millman, St’efan J. van der Walt, Ralf Gommers, Pauli Virtanen, David Cournapeau, Eric Wieser, Julian Taylor, Sebastian Berg, Nathaniel J. Smith, Robert Kern, Matti Picus, Stephan Hoyer, Marten H. van Kerkwijk, Matthew Brett, Allan Haldane, Jaime Fern’andez del R’io, Mark Wiebe, Pearu Peterson, Pierre G’erard-Marchant, Kevin Sheppard, Tyler Reddy, Warren Weckesser, Hameer Abbasi, Christoph Gohlke, and Travis E. Oliphant. Array programming with NumPy. *Nature*, 585(7825):357–362, September 2020.
- [8] Laurence James. Automatic requirements specification update processing from a requirements management tool perspective. In *Proceedings of the International Conference and Workshop on Engineering of Computer-Based Systems*, pages 2–9. IEEE, 1997.
- [9] Guillaume Lemaître, Fernando Nogueira, and Christos K. Aridas. Imbalanced-learn: A python toolbox to tackle the curse of imbalanced datasets in machine learning. *Journal of Machine Learning Research*, 18(17):1–5, 2017.
- [10] Patrick Mäder and Alexander Egyed. Do developers benefit from requirements traceability when evolving and maintaining a software system? *Empirical Software Engineering*, 20(2):413–441, 2015.
- [11] Chris Mills, Javier Escobar-Avila, Aditya Bhattacharya, Grigoriy Kondyukov, Shayok Chakraborty, and Sonia Haiduc. Tracing with less data: Active learning for classification-based traceability link recovery. In *2019 IEEE International Conference on Software Maintenance and Evolution (ICSME)*, pages 103–113. IEEE, 2019.
- [12] Chris Mills, Javier Escobar-Avila, and Sonia Haiduc. Automatic traceability maintenance via machine learning classification. In *2018 IEEE International Conference on Software Maintenance and Evolution (ICSME)*, pages 369–380. IEEE, 2018.
- [13] Patrick Rempel and Patrick Mader. Preventing defects: The impact of requirements traceability completeness on software quality. *IEEE Transactions on Software Engineering*, 2016.
- [14] Klaus Weidenhaupt, Klaus Pohl, Matthias Jarke, and Peter Haumer. Scenarios in system development: current practice. *IEEE Software*, 15(2):34–45, 1998.